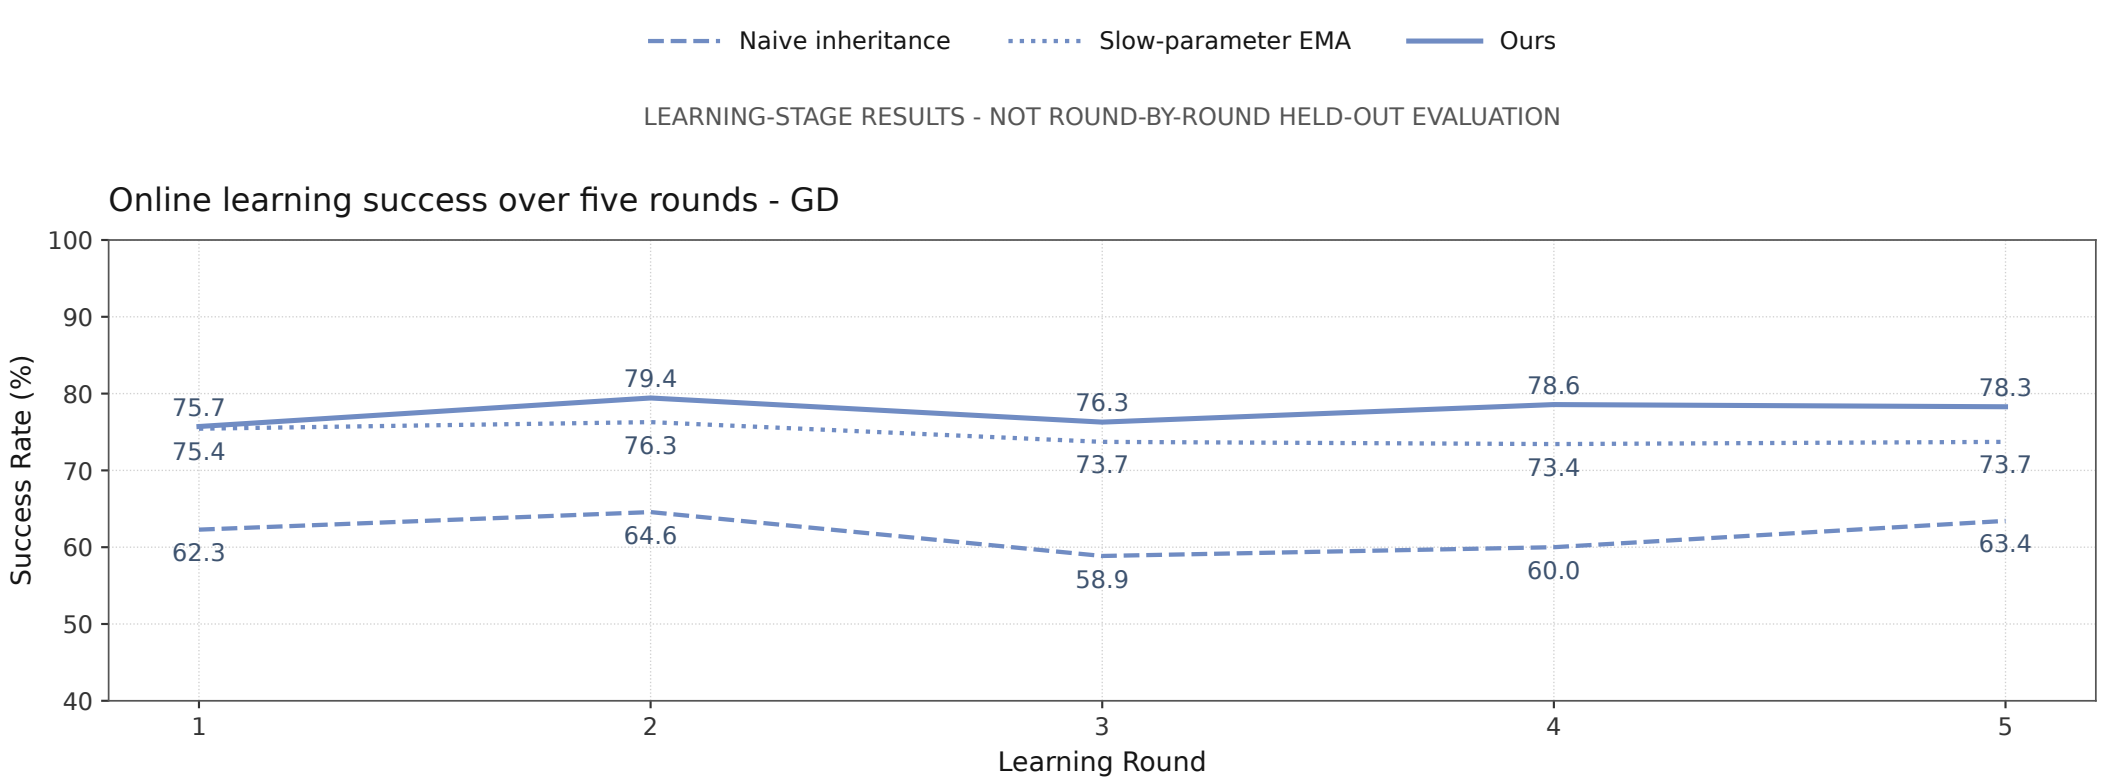

350 learning episodes per round and method; one chain. No error bars from training repeats.

Learning seeds change by round (100, 200, 300, 1000, 1100); values are not fixed-test learning curves.

AdaJEPA: no comparable five-round learning log. All four methods appear in held-out Figures 1-3.